Topic 3: What Is RAG?

Simple RAG Project - Requirements Guide

Break Retrieval-Augmented Generation into retrieval, augmentation, and generation.

What this topic requires

Retrieval finds useful source chunks; augmentation adds them to the prompt; generation asks the LLM to produce the response.

Required understanding and workflow

1. Convert the question into a searchable representation.

2. Find relevant document chunks.

3. Add those chunks to the prompt.

4. Call the LLM.

5. Return a grounded answer.

Workflow

```
Question -> embedding -> search -> chunks -> augmented prompt -> LLM -> answer
```

Evidence rule

Describe implementation only after repository inspection. Missing features must be labeled NOT IMPLEMENTED IN THIS PROJECT. Recommendations must be labeled PRODUCTION RECOMMENDATION - NOT CURRENTLY IMPLEMENTED.