Topic 10: Embeddings

Simple RAG Project - Requirements Guide

Explain how text is converted into vectors so semantically similar text can be found.

What this topic requires

Document the exact embedding model, provider, dimensions only when identifiable, input/output behavior, cost, and local-versus-cloud trade-off.

Required understanding and workflow

1. Identify the embedding integration.

2. Record model and provider.

3. Verify vector dimension if exposed.

4. Explain query and document embeddings.

5. Explain semantic similarity without inventing model details.

Workflow

```
Text -> embedding model -> vector [numbers...]
```

Evidence rule

Describe implementation only after repository inspection. Missing features must be labeled NOT IMPLEMENTED IN THIS PROJECT. Recommendations must be labeled PRODUCTION RECOMMENDATION - NOT CURRENTLY IMPLEMENTED.